



# **Database Systems**

## **Lecture 11: Advanced SQL (part 1)**

**Dr. Momtazi**  
**[momtazi@aut.ac.ir](mailto:momtazi@aut.ac.ir)**

based on the slides of the course book



# Outline

- **Accessing SQL From a Programming Language**
- Functions
- Triggers
- Advanced Aggregation Features
- OLAP



# Accessing SQL From a Programming Language

- API (application-program interface) for a program to interact with a database server
  
- Application makes calls to
  - Connect with the database server
  - Send SQL commands to the database server
  - Fetch tuples of result one-by-one into program variables



# Accessing SQL From a Programming Language

- Possible approaches:
  - Dynamic SQL
    - ▶ JDBC (Java Database Connectivity) works with Java
    - ▶ ODBC (Open Database Connectivity) works with C, C++, C#, and Visual Basic.
      - Other API's such as ADO.NET sit on top of ODBC
  - Embedded SQL



# JDBC

- **JDBC** is a Java API for communicating with database systems supporting SQL.
- JDBC supports a variety of features for querying and updating data, and for retrieving query results.
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes.



# JDBC

- Model for communicating with the database:
  - Open a connection
  - Create a “statement” object
  - Execute queries using the Statement object to send queries and fetch results
  - Exception mechanism to handle errors



# JDBC Example

```
public static void JDBCexample(String userid, String passwd)
{
    try
    {
        Class.forName ("oracle.jdbc.driver.OracleDriver");
        Connection conn = DriverManager.getConnection(
            "jdbc:oracle:thin:@db.yale.edu:1521:univdb",
            userid, passwd);
        Statement stmt = conn.createStatement();
        try {
            stmt.executeUpdate(
                "insert into instructor values('77987', 'Kim', 'Physics', 98000)");
        } catch (SQLException sqle)
        {
            System.out.println("Could not insert tuple. " + sqle);
        }
        ResultSet rset = stmt.executeQuery(
            "select dept_name, avg (salary) "+
            " from instructor "+
            " group by dept_name");
        while (rset.next()) {
            System.out.println(rset.getString("dept_name") + " " +
                               rset.getFloat(2));
        }
        stmt.close();
        conn.close();
    }
    catch (Exception sqle)
    {
        System.out.println("Exception : " + sqle);
    }
}
```

Apago PDF Enhancer



# JDBC Example

```
public static void JDBCexample(String userid, String passwd)
{
    try
    {
        Class.forName ("oracle.jdbc.driver.OracleDriver");
        Connection conn = DriverManager.getConnection(
            "jdbc:oracle:thin:@db.yale.edu:1521:univdb",
            userid, passwd);
        Statement stmt = conn.createStatement();
        try {
            stmt.executeUpdate(
                "insert into instructor values('77987', 'Kim', 'Physics', 98000)");
        } catch (SQLException sqle)
        {
            System.out.println("Could not insert tuple. " + sqle);
        }
    }
}
```





# JDBC Example

```
ResultSet rset = stmt.executeQuery(
    "select dept_name, avg (salary) "+
    " from instructor "+
    " group by dept_name");
while (rset.next()) {
    System.out.println(rset.getString("dept_name") + " " +
        rset.getFloat(2));
}
stmt.close();
conn.close();
}
catch (Exception sqle)
{
    System.out.println("Exception : " + sqle);
}
}
```

Apago PDF Enhancer



# Database Connection

- Each database product that supports JDBC provides a JDBC driver that must be dynamically loaded in order to access the database from Java.
  - This is done by invoking *Class.forName* with one argument specifying a concrete class implementing the `java.sql.Driver` interface
- Connecting to the Database: A connection is opened using the `getConnection` method of the `DriverManager` class (within `java.sql`) using 3 parameters:
  - a string that specifies the URL, or machine name, where the server runs
  - a database user identifier, which is a string
  - a password, which is also a string.
    - ▶ Note: the need to specify a password within the JDBC code presents a security risk if an unauthorized person accesses your Java code.



# Shipping SQL Statements

- Methods for executing a statement:
  - *executeQuery*
    - ▶ When the SQL statement is a query
    - ▶ It returns a result set
  - *executeUpdate*
    - ▶ When the SQL statement is nonquery (DDL or DML)
      - Update
      - Insert
      - Delete
      - Create table
      - ...
    - ▶ It returns an integer giving the number of tuples inserted, updated, or deleted.
    - ▶ For DDL statements, the return value is zero.



# Retrieving the Results of a Query

- Retrieving the set of tuples in the result into a `ResultSet` object
- Fetching the results one tuple at a time
- Using the *next* method on the result set to test whether there remains at least one unfetched tuple in the result set and if so, fetches it.
- Attributes from the fetched tuple are retrieved using various methods whose names begin with *get*
  - *getString*: can retrieve any of the basic SQL data types
  - *getFloat*
- Possible argument to the *get* methods
  - The attribute name specified as a string
  - An integer indicating the position of the desired attribute within the tuple



# Database Connection

- The statement and connection are both closed at the end of the Java program.
- It is important to close the connection because there is a limit imposed on the number of connections to the database
- Unclosed connections may cause that limit to be exceeded.
- If this happens, the application cannot open any more connections to the database.



# Prepared Statements

- Creating a prepared statement in which some values are replaced by “?”
- Specifying that actual values will be provided later
- Compiling the query by the database system when it is prepared
- Reusing the previously compiled form of the query and apply the new values whenever the query is executed
  - (with new values to replace the “?”s),



# Prepared Statements

```
PreparedStatement pStmt = conn.prepareStatement(
    "insert into instructor values(?,?,?,?)");
pStmt.setString(1, "88877");
pStmt.setString(2, "Perry");
pStmt.setString(3, "Finance");
pStmt.setInt(4, 125000);
pStmt.executeUpdate();
pStmt.setString(1, "88878");
pStmt.executeUpdate();
```

**Figure 5.2** Prepared statements in JDBC code.

- insert into instructor values ("88877", "Perry", "Finance", 125000)
- insert into instructor values ("88878", "Perry", "Finance", 125000)



# Prepared Statements

- Prepared statements allow for more efficient execution
  - Where the same query can be compiled once and then run multiple times with different parameter values
  - Whenever a user-entered value is used, even if the query is to be run only once





# Metadata Features

- Capturing metadata about
  - Database
  - ResultSet (relations)

```
ResultSetMetaData rsmd = rs.getMetaData();
for(int i = 1; i <= rsmd.getColumnCount(); i++) {
    System.out.println(rsmd.getColumnName(i));
    System.out.println(rsmd.getColumnTypeName(i));
}
```



# ODBC

- Open DataBase Connectivity (ODBC) standard
  - standard for application program to communicate with a database server.
  - application program interface (API) to
    - ▶ open a connection with a database,
    - ▶ send queries and updates,
    - ▶ get back results.



# Embedded SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as C, C++, Java, Fortran, and PL/1,
- A language to which SQL queries are embedded is referred to as a **host language**
- The SQL structures permitted in the host language comprise *embedded SQL*
- An embedded SQL program must be processed by a special preprocessor prior to compilation.
- The preprocessor replaces embedded SQL requests with host-language declarations and procedure calls that allow runtime execution of the database accesses.
- Then, the resulting program is compiled by the host-language compiler.



# Embedded SQL

- **EXEC SQL** statement is used to identify embedded SQL request to the preprocessor

EXEC SQL <embedded SQL statement >;

- Note: this varies by language:
  - In some languages, like COBOL, the semicolon is replaced with END-EXEC
  - In Java embedding uses  
# SQL { .... };



# Database Connection

- Before executing any SQL statements, the program must first connect to the database. This is done using:

EXEC-SQL **connect to** *server* **user** *user-name* **using** *password*;

- Here, *server* identifies the server to which a connection is to be established.



# Variables

- Variables of the host language can be used within embedded SQL statements.
- They are preceded by a colon (:) to distinguish from SQL variables (e.g., *:credit\_amount* )
- Variables used as above must be declared within DECLARE section, as illustrated below. The syntax for declaring the variables, however, follows the usual host language syntax.

```
EXEC-SQL BEGIN DECLARE SECTION;
```

```
    int credit_amount ;
```

```
EXEC-SQL END DECLARE SECTION;
```



# SQL Query

- To write an embedded SQL query, we use the following statement:

**declare c cursor for <SQL query>**

- The variable *c* is used to identify the query

- Example:

**EXEC SQL**

**declare c cursor for**

**select *ID, name***

**from *student***

**where *tot\_cred* > *:credit\_amount***

**END\_EXEC**



# SQL Query

- The open statement is then used to evaluate the query
- The **open** statement for our example is as follows:

**EXEC SQL open c ;**

- This statement causes the database system to execute the query and to save the results within a temporary relation.
- The query uses the value of the host-language variable *credit-amount* at the time the **open** statement is executed.





# SQL Query

## ■ The fetch statement

- Placing the values of one tuple in the query result into host language variables
- Requiring one host-language variable for each attribute of the result relation;

e.g., we need one variable to hold the ID value (si) and another to hold the name value (sn) which have been declared within a DECLARE section

**EXEC SQL**

**fetch c into :si, :sn**

**END\_EXEC**

Repeated calls to fetch get successive tuples in the query result



# SQL Query

- The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

**EXEC SQL close c ;**



# Updates Through Embedded SQL

- Embedded SQL expressions for database modification (**update**, **insert**, and **delete**)
- Example for updating tuples fetched by cursor by declaring that the cursor is for update

## EXEC SQL

```
declare c cursor for  
select *  
from instructor  
where dept_name = 'Music'  
for update;
```



# Updates Through Embedded SQL

- Iterating through the tuples by performing **fetch** operations on the cursor
- Executing the following statement

## EXEC SQL

```
update instructor  
  set salary = salary + 1000  
  where current of c;
```



# Outline

- Accessing SQL From a Programming Language
- **Functions**
- Triggers
- Advanced Aggregation Features
- OLAP



# Functions and Procedures

- SQL:1999 supports functions and procedures
  - Functions/procedures can be written in SQL itself, or in an external programming language (e.g., C, Java).
  - Some database systems support **table-valued functions**, which can return a relation as a result.
- SQL:1999 also supports a rich set of imperative constructs, including
  - Loops, if-then-else, assignment
- Many databases have proprietary procedural extensions to SQL that differ from SQL:1999.



# SQL Functions

- Define a function that, given the name of a department, returns the count of the number of instructors in that department.

```
create function dept_count (dept_name varchar(20))  
  returns integer  
  begin  
    declare d_count integer;  
    select count ( * ) into d_count  
    from instructor  
    where instructor.dept_name = dept_name  
    return d_count;  
end
```



# SQL Functions

- The function *dept\_count* can be used in a query that returns names and budget of all departments with more that 12 instructors.

```
select dept_name, budget  
from department  
where dept_count (dept_name ) > 12
```





# SQL Functions

- Compound statement: **begin ... end**
  - May contain multiple SQL statements between **begin** and **end**.
- **returns** -- indicates the variable-type that is returned (e.g., integer)
- **return** -- specifies the values that are to be returned as result of invoking the function
- SQL function are in fact **parameterized views** that generalize the regular notion of views by allowing parameters.



# Table Functions

- SQL:2003 added functions that return a relation as a result
- Example: Return all instructors in a given department

**create function** *instructor\_of* (*dept\_name* **char**(20))

**returns table** (

*ID* **varchar**(5),  
*name* **varchar**(20),  
*dept\_name* **varchar**(20),  
*salary* **numeric**(8,2))

**return table**

(**select** *ID*, *name*, *dept\_name*, *salary*  
**from** *instructor*  
**where** *instructor.dept\_name* = *instructor\_of.dept\_name*)

- Usage

**select** \*  
**from table** (*instructor\_of* ('Music'))



Questions?