



# **Database Systems**

## **Lecture 12: Advanced SQL (part 2)**

**Dr. Momtazi**  
**[momtazi@aut.ac.ir](mailto:momtazi@aut.ac.ir)**

based on the slides of the course book



# Outline

- Accessing SQL From a Programming Language
- Functions
- **Triggers**
- Advanced Aggregation Features
- OLAP



# Triggers

- A **trigger** is a statement that is executed automatically by the system as a side effect of a modification to the database.
- To design a trigger mechanism, we must:
  - Specify the conditions under which the trigger is to be executed.
  - Specify the actions to be taken when the trigger executes.
- Triggers introduced to SQL standard in SQL:1999, but supported even earlier using non-standard syntax by most databases.
  - Syntax illustrated here may not work exactly on your database system; check the system manuals



# Triggering Events and Actions in SQL

- Triggering event can be **insert**, **delete** or **update**
- Triggers on update can be restricted to specific attributes
  - For example, **after update of *takes on grade***
- Values of attributes before and after an update can be referenced
  - **referencing old row as** : for deletes and updates
  - **referencing new row as** : for inserts and updates
- Triggers can be activated before an event, which can serve as extra constraints. For example, convert blank grades to null.



# Trigger Example: Using set Statement

```
create trigger setnull before update on takes  
referencing new row as nrow  
for each row  
when (nrow.grade = ' ')  
begin atomic  
    set nrow.grade = null;  
end;
```



# Trigger Example: to Maintain Referential Integrity

```
create trigger timeslot_check1 after insert on section
referencing new row as nrow
for each row
when (nrow.time_slot_id not in (
    select time_slot_id
    from time_slot)) /* time_slot_id not present in time_slot */
begin
    rollback
end;

create trigger timeslot_check2 after delete on timeslot
referencing old row as orow
for each row
when (orow.time_slot_id not in (
    select time_slot_id
    from time_slot) /* last tuple for time_slot_id deleted from time_slot */
and orow.time_slot_id in (
    select time_slot_id
    from section)) /* and time_slot_id still referenced from section*/
begin
    rollback
end;
```

Apache PDF Enhancer



# Trigger Example: to Maintain `credits_earned` value

```
create trigger credits_earned after update of takes on (grade)
referencing new row as nrow
referencing old row as orow
for each row
when nrow.grade <> 'F' and nrow.grade is not null
  and (orow.grade = 'F' or orow.grade is null)
begin atomic
  update student
  set tot_cred = tot_cred +
    (select credits
     from course
     where course.course_id = nrow.course_id)
  where student.id = nrow.id;
end;
```

Figure 5.9 Using a trigger to maintain *credits\_earned* values.



# Disabling Triggers

- Triggers can be disabled or enabled;
  - by default they are enabled when they are created.
- Triggers can be disabled by:  
**alter trigger *trigger\_name* disable**  
**disable trigger *trigger\_name***
- A trigger that has been disabled can be enabled again.
- A trigger can instead be dropped, which removes it permanently, by:  
**drop trigger *trigger\_name***





# When Not To Use Triggers

- Triggers were used earlier for tasks such as
  - Maintaining summary data (e.g., total salary of each department)
  - Replicating databases by recording changes to special relations (called **change** or **delta** relations) and having a separate process that applies the changes over to a replica
- There are better ways of doing these now:
  - Databases today provide built in materialized view facilities to maintain summary data
  - Databases provide built-in support for replication
- Encapsulation facilities can be used instead of triggers in many cases
  - Define methods to update fields
  - Carry out actions as part of the update methods instead of through a trigger



# When Not To Use Triggers

- Risk of unintended execution of triggers, for example, when
  - Loading data from a backup copy
  - Replicating updates at a remote site
  - Trigger execution can be disabled before such actions.
- Other risks with triggers:
  - Error leading to failure of critical transactions that set off the trigger
  - Cascading execution



# Outline

- Accessing SQL From a Programming Language
- Functions
- Triggers
- **Advanced Aggregation Features**
- OLAP



# Ranking

- Suppose we are given a relation  
*student\_grades*(*ID*, *GPA*)  
giving the grade-point average of each student
- Goal: finding the rank of each student.
- Ranking can be done using basic SQL aggregation, but resultant query is very inefficient

```
select ID, (1 + (select count(*)  
                    from student_grades B  
                    where B.GPA > A.GPA)) as s_rank  
from student_grades A  
order by s_rank;
```



# Ranking

- Ranking is done in conjunction with an order by specification.

```
select ID, rank() over (order by GPA desc) as s_rank  
from student_grades  
order by s_rank
```

- NOTE: the extra **order by** clause is needed to get them in sorted order



# Ranking

- Two possible approaches for ranking
  - Leaving gaps: (default)  
e.g. if 2 students have the same top GPA, both have rank 1, and the next rank is 3
  - Without gaps: (using **dense\_rank**)  
so next dense rank would be 2



# Ranking with Partitions

- Ranking can be done within partition of the data.
- “Find the rank of students within each department.”

```
select ID, dept_name,  
        rank () over (partition by dept_name order by GPA desc)  
        as dept_rank  
from dept_grades  
order by dept_name, dept_rank;
```

- Ranking is done *after* applying **group by** clause/aggregation



# Top $n$ Items

- Can be used to find top- $n$  results
  - More general than the **limit**  $n$  clause supported by many databases, since it allows top- $n$  within each partition





# Other Ranking Functions

## ■ **percent\_rank**

- within partition, if partitioning is done
- Having  $n$  tuples in a partition and the rank  $r$  for the target tuple
- $percent\_rank = (r-1)/(n-1)$



# Other Ranking Functions

- **cume\_dist** (cumulative distribution)
  - fraction of tuples with preceding values
  - Having  $n$  tuples in a partition and  $p$  tuples with ordering value preceding or equal to the the target tuple
  - $cume\_dist = p/n$



# Other Ranking Functions

## ■ **row\_number**

- Sorting the rows and giving each row a unique number corresponding to its position in the sort order
- Non-deterministic in presence of duplicates  
(different rows with the same ordering value would get different random row numbers)



# Other Ranking Functions

## ■ **ntile:**

- For a given constant  $n$ , the ranking the function  $ntile(n)$  takes the tuples in each partition in the specified order, and divides them into  $n$  buckets with equal numbers of tuples.
- E.g.,

```
select ID, ntile(4) over (order by GPA desc) as quartile  
from student_grades;
```



# Other Ranking Functions

- SQL:1999 permits the user to specify **nulls first** or **nulls last**

```
select ID,  
        rank ( ) over (order by GPA desc nulls last) as s_rank  
from student_grades
```



# Outline

- Accessing SQL From a Programming Language
- Functions
- Triggers
- Advanced Aggregation Features
- **OLAP**



# Data Analysis and OLAP

## ■ Online Analytical Processing (OLAP)

- Interactive analysis of data, allowing data to be summarized and viewed in different ways in an online fashion (with negligible delay)
- Is used for **multidimensional data** (data that can be modeled as dimension attributes and measure attributes)
  - **Measure attributes**
    - ▶ measure some value
    - ▶ can be aggregated upon
    - ▶ e.g., the attribute *number* of the *sales* relation
  - **Dimension attributes**
    - ▶ define the dimensions on which measure attributes (or aggregates there of) are viewed
    - ▶ e.g., attributes *item\_name*, *color*, and *size* of the *sales* relation



## Example sales relation

| <i>item_name</i> | <i>color</i> | <i>clothes_size</i> | <i>quantity</i> |
|------------------|--------------|---------------------|-----------------|
| skirt            | dark         | small               | 2               |
| skirt            | dark         | medium              | 5               |
| skirt            | dark         | large               | 1               |
| skirt            | pastel       | small               | 11              |
| skirt            | pastel       | medium              | 9               |
| skirt            | pastel       | large               | 15              |
| skirt            | white        | small               | 2               |
| skirt            | white        | medium              | 5               |
| skirt            | white        | large               | 3               |
| dress            | dark         | small               | 2               |
| dress            | dark         | medium              | 6               |
| dress            | dark         | large               | 12              |
| dress            | pastel       | small               | 4               |
| dress            | pastel       | medium              | 3               |
| dress            | pastel       | large               | 3               |
| dress            | white        | small               | 2               |
| dress            | white        | medium              | 3               |
| dress            | white        | large               | 0               |
| shirt            | dark         | small               | 2               |
| shirt            | dark         | medium              | 4               |
| ...              | ...          | ...                 | ...             |
| ...              | ...          | ...                 | ...             |





# Cross Tabulation of sales by *item\_name* and color

*clothes\_size* **all**


|                  |       | <i>color</i> |        |       |       |
|------------------|-------|--------------|--------|-------|-------|
|                  |       | dark         | pastel | white | total |
| <i>item_name</i> | skirt | 8            | 35     | 10    | 53    |
|                  | dress | 20           | 10     | 5     | 35    |
|                  | shirt | 14           | 7      | 28    | 49    |
|                  | pants | 20           | 2      | 5     | 27    |
|                  | total | 62           | 54     | 48    | 164   |

- The table above is an example of a **cross-tabulation** (**cross-tab**), also referred to as a **pivot-table**.
  - Values for one of the dimension attributes form the row headers
  - Values for another dimension attribute form the column headers
  - Other dimension attributes are listed on top
  - Values in individual cells are (aggregates of) the values of the dimension attributes that specify the cell.



# Data Cube

- A **data cube** is a multidimensional generalization of a cross-tab
- Can have  $n$  dimensions; we show 3 below
- Cross-tabs can be used as views on a data cube





# Online Analytical Processing Operations

- **Pivoting:** changing the dimensions used in a cross-tab is called
- **Slicing:** creating a cross-tab for fixed values only
  - Sometimes called **dicing**, particularly when values for multiple dimensions are fixed.



# Cross Tabulation With Hierarchy

- The following table can be achieved using natural join with another table specifying item\_names and category

*clothes\_size:* **all**

| <i>category</i> | <i>item_name</i> | <i>color</i> |        |       |       |     |
|-----------------|------------------|--------------|--------|-------|-------|-----|
|                 |                  | dark         | pastel | white | total |     |
| womenswear      | skirt            | 8            | 8      | 10    | 53    | 88  |
|                 | dress            | 20           | 20     | 5     | 35    |     |
|                 | subtotal         | 28           | 28     | 15    |       |     |
| menswear        | pants            | 14           | 14     | 28    | 49    | 76  |
|                 | shirt            | 20           | 20     | 5     | 27    |     |
|                 | subtotal         | 34           | 34     | 33    |       |     |
| total           |                  | 62           | 62     | 48    |       | 164 |



# Cross Tabulation With Hierarchy

- Cross-tabs can be easily extended to deal with hierarchies
- Can drill down or roll up on a hierarchy
  - roll up: moving from finer granularity to coarser-granularity (by the mean of aggregation)
  - drill down: moving from coarser-granularity to finer granularity (must be generated from original data)



# Relational Representation of Cross-tabs

- Cross-tabs can be represented as relations
  - We use the value **all** is used to represent aggregates.
  - The SQL standard actually uses null values in place of **all** despite confusion with regular null values.

| <i>item_name</i> | <i>color</i> | <i>clothes_size</i> | <i>quantity</i> |
|------------------|--------------|---------------------|-----------------|
| skirt            | dark         | <b>all</b>          | 8               |
| skirt            | pastel       | <b>all</b>          | 35              |
| skirt            | white        | <b>all</b>          | 10              |
| skirt            | <b>all</b>   | <b>all</b>          | 53              |
| dress            | dark         | <b>all</b>          | 20              |
| dress            | pastel       | <b>all</b>          | 10              |
| dress            | white        | <b>all</b>          | 5               |
| dress            | <b>all</b>   | <b>all</b>          | 35              |
| shirt            | dark         | <b>all</b>          | 14              |
| shirt            | pastel       | <b>all</b>          | 7               |
| shirt            | White        | <b>all</b>          | 28              |
| shirt            | <b>all</b>   | <b>all</b>          | 49              |
| pant             | dark         | <b>all</b>          | 20              |
| pant             | pastel       | <b>all</b>          | 2               |
| pant             | white        | <b>all</b>          | 5               |
| pant             | <b>all</b>   | <b>all</b>          | 27              |
| <b>all</b>       | dark         | <b>all</b>          | 62              |
| <b>all</b>       | pastel       | <b>all</b>          | 54              |
| <b>all</b>       | white        | <b>all</b>          | 48              |
| <b>all</b>       | <b>all</b>   | <b>all</b>          | 164             |



# Extended Aggregation to Support OLAP

- The **cube** operation computes union of **group by**'s on every subset of the specified attributes
- Example relation for this section  
*sales(item\_name, color, clothes\_size, quantity)*
- E.g. consider the query

```
select item_name, color, size, sum(number)  
from sales  
group by cube(item_name, color, size)
```

This computes the union of eight different groupings of the *sales* relation:

```
{ (item_name, color, size), (item_name, color),  
  (item_name, size),      (color, size),  
  (item_name),           (color),  
  (size),                ( ) }
```

where ( ) denotes an empty **group by** list.

- For each grouping, the result contains the null value for attributes not present in the grouping.



# Extended Aggregation

- The **rollup** construct generates union on every prefix of specified list of attributes

- E.g.,

```
select item_name, color, size, sum(number)  
from sales  
group by rollup(item_name, color, size)
```

Generates union of four groupings:

```
{ (item_name, color, size), (item_name, color), (item_name), ( ) }
```





# Extended Aggregation

- Rollup can be used to generate aggregates at multiple levels of a hierarchy.
- E.g., suppose table *itemcategory*(*item\_name*, *category*) gives the category of each item. Then

```
select category, item_name, sum(number)
from sales, itemcategory
where sales.item_name = itemcategory.item_name
group by rollup(category, item_name)
```

would give a hierarchical summary by *item\_name* and by *category*.



# Extended Aggregation

- Multiple rollups or cubes can be used in a single group by clause
  - Each generates set of group by lists, cross product of sets gives overall set of group by lists

■ E.g.,

```
select item_name, color, size, sum(number)  
from sales  
group by rollup(item_name), rollup(color, size)
```

generates the groupings

$$\{item\_name, ()\} \times \{(color, size), (color), ()\}$$
$$= \{ (item\_name, color, size), (item\_name, color), (item\_name), (color, size), (color), ( ) \}$$



# OLAP Implementation

- The earliest OLAP systems used multidimensional arrays in memory to store data cubes, and are referred to as **multidimensional OLAP (MOLAP)** systems.
- OLAP implementations using only relational database features are called **relational OLAP (ROLAP)** systems
- Hybrid systems, which store some summaries in memory and store the base data and other summaries in a relational database, are called **hybrid OLAP (HOLAP)** systems.



# OLAP Implementation

- Early OLAP systems precomputed *all* possible aggregates in order to provide online response
  - Space and time requirements for doing so can be very high
    - ▶  $2^n$  combinations of **group by**
- Several optimizations available for computing multiple aggregates
  - It suffices to precompute some aggregates, and compute others on demand from one of the precomputed aggregates
    - ▶ Can compute aggregate on *(item\_name, color)* from an aggregate on *(item\_name, color, size)*
      - is cheaper than computing it from scratch



Questions?